

Predicting Customer Churn for a Telecommunication Company



School of Computer Science and Engineering

Lovely Professional University

Phagwara, Punjab

NAME: Shyam

REGISTRATION NO.: 12111001

1.Introduction:

- Project Objective:

Predict customer churn using machine learning techniques.

- Data Source:

'Churn.csv'

- Tools and Libraries Used:

Pandas, NumPy, Matplotlib, Seaborn, Scikit-learn, XGBoost

2. Data Loading and Preprocessing

- 2.1 Import Libraries:

- **pandas (pd)**

- Uses:

1. **Data Loading and Manipulation:** Used to load data from CSV files (e.g., `pd.read_csv`) and manipulate it using dataframes for cleaning, transformation, and analysis (e.g., selecting rows/columns, filtering data).
2. **Data Analysis:** Provides functions for descriptive statistics (e.g., mean, median, standard deviation), time series analysis (e.g., resampling data), and creating data visualizations (e.g., plotting dataframes).

- **numpy (np)**

- Uses:

1. **Numerical Computations:** Offers efficient arrays for numerical operations (e.g., vectorized calculations, matrix operations) like element-wise addition, multiplication, and linear algebra functions.
2. **Data Analysis:** Provides functions for array manipulation (e.g., reshaping, indexing), random number generation (e.g., generating random samples from various distributions), and linear algebra operations (e.g., solving systems of equations).

- **seaborn (sns)**

- Uses:

1. **Statistical Visualization:** Built on top of matplotlib, it provides high-level functions for creating informative and aesthetically pleasing statistical graphics like boxplots, violin plots, and jointplots to visualize relationships between variables.
2. **Distribution Plots:** Offers functions for creating histograms, kernel density plots, and cumulative distribution functions to explore the distribution of data.

- **matplotlib.pyplot (plt)**

- Uses:

1. **Basic Plotting:** Provides functions for creating basic plots like line plots, scatter plots, and bar charts to visualize data trends and relationships.
 2. **Customization:** Allows for customization of plot elements like labels, titles, colors, and legends for clear and informative data representation.
- **sklearn.linear_model (LogisticRegression)**
 - Uses:
 1. **Classification:** Implements the Logistic Regression algorithm for binary classification problems, predicting the probability of an outcome based on features (e.g., churn prediction).
 2. **Regression Analysis:** Can also be used for linear regression tasks to model the relationship between a continuous dependent variable and one or more independent variables.
 - **sklearn.metrics (accuracy_score, confusion_matrix, classification_report)**
 - Uses:
 1. **Model Evaluation:** Provides functions like `accuracy_score` to measure overall model performance, `confusion_matrix` to visualize the distribution of true vs. predicted labels, and `classification_report` for detailed metrics like precision, recall, and F1-score.
 - **sklearn.preprocessing (LabelEncoder)**
 - Uses:
 1. **Categorical Encoding:** Used for converting categorical data (e.g., text labels) into numerical values suitable for machine learning algorithms. `LabelEncoder` assigns a unique integer to each category.
 2. **Data Scaling:** Other functions like `StandardScaler` can be used to scale numerical features to a common range for improved model performance.
 - **imblearn.over_sampling (SMOTE)**
 - Uses:
 1. **Data Balancing:** Addresses class imbalance problems in datasets, where one class has significantly fewer samples than others. SMOTE (Synthetic Minority Oversampling Technique) creates synthetic samples for the minority class.
 2. **Under-Sampling:** Other techniques like random under-sampling can also be used to reduce the majority class size and achieve a more balanced dataset.
 - **collections (Counter)**
 - Uses:
 1. **Frequency Analysis:** Provides the `Counter` class to count the occurrences of elements in an iterable object (e.g., list, dictionary). Useful for exploring the distribution of categorical data or imbalanced classes.

2. **Set Operations:** Offers functions for working with sets, such as finding the union, intersection, and difference between sets, which can be helpful for data manipulation tasks.

- **sklearn.model_selection (train_test_split)**

- Uses:

1. **Data Splitting:** Splits a dataset into training and testing sets for model training and evaluation. train_test_split allows for specifying the desired test size ratio and randomization for robust model assessment.

- **sklearn.tree (DecisionTreeClassifier)**

- Uses:

1. **Classification:** Implements the Decision Tree algorithm for classification tasks. It builds a tree-like structure based on feature values to predict class labels.
2. **Feature Importance:** Can be used to understand the relative importance of features in the model by analyzing the splits made at each node in the tree.

- **sklearn.ensemble (RandomForestClassifier, GradientBoostingClassifier)**

- Uses:

1. **Ensemble Learning:** Both Random Forest and Gradient Boosting are ensemble methods that combine multiple weak learners (e.g., decision trees) to create a stronger model.
2. **Model Regularization:** These techniques help to reduce

- **sklearn.neighbors (KNeighborsClassifier)**

- Uses:

1. **Classification:** Used for classifying data points based on the k nearest neighbors in the training data. It predicts the class label that is most frequent among the k nearest neighbors for a new data point.

- **sklearn.svm (SVC)**

- Uses:

1. **Classification:** Implements Support Vector Machines (SVM) for classification tasks. It finds a hyperplane that best separates the data points of different classes with the maximum margin.

- **xgboost (XGBClassifier)**

- Uses:

1. **Classification (Advanced):** XGBoost is a powerful machine learning library for Gradient Boosting, a technique that builds an ensemble of decision trees sequentially. It's often used for complex classification problems due to its accuracy and efficiency.

- **plotly.express (px)**

- Uses:

1. **Interactive Data Visualization:** Provides a high-level interface for creating interactive visualizations like scatter plots, bar charts, and heatmaps using plotly.

Useful for exploring relationships between variables and understanding model predictions.

- **statsmodels.api (sm)**

-Uses:

1. **Statistical Modeling:** Offers a comprehensive library for statistical modeling and econometrics. While it can be used for various tasks, a common use case in machine learning is for fitting statistical models like linear regression and analyzing their properties.

2.2 Load Data:

- Data File: 'Churn.csv'

- Dataframe: churn_data

```
) churn_data=pd.read_csv('/content/sample_data/Churn.csv')
```

2.3 Data Cleaning:

Missing Values:

Strategy: Converted empty strings in 'totalcharges' to the mean value.

```
[304] churn_data['totalcharges'].describe()
```



```
count    7032.000000
mean     2283.300441
std      2266.771362
min       18.800000
25%       401.450000
50%      1397.475000
75%      3794.737500
max      8684.800000
Name: totalcharges, dtype: float64
```


```
[305] churn_data['totalcharges'].fillna(value=2283.300441,inplace=True)
```

Outliers:

- Identification: Used boxplots.

```
columns=['tenure','monthlycharges','totalcharges']  
for i in columns:  
    sns.boxplot(churn_data[i])  
    plt.show()
```





- Decision: Kept outliers as they were not extreme.

Duplicates:

- Action: duplicates are not present.

```
churn_data.duplicated().sum()
```

0

Irrelevant Columns:

- Action: Removed 'customerID'.

customerid is not correlated to the target variable.so I remove the customerid column.

```
[311] churn_data.drop('customerid', axis=1, inplace=True)
      churn_data.head()
```

Categorical Data:

_Change abnormal categorical values:

```
churn_data['multiplelines'] = churn_data['multiplelines'].replace('No phone service', 'No')
```

```
data_columns = ['internetservice', 'onlinesecurity', 'onlinebackup', 'deviceprotection', 'techsupport',
                'streamingtv', 'streamingmovies']

for column in data_columns:
    churn_data[column] = churn_data[column].replace('No internet service', 'No')
```

- Strategy: Converted categorical data to numerical values using LabelEncoder.

```
categorical = list(churn_data.select_dtypes(include=['object']).columns)
categorical
```


```
from sklearn.preprocessing import LabelEncoder
LB=LabelEncoder()
for i in categorical:
    churn_data[i]=LB.fit_transform(churn_data[i])
```

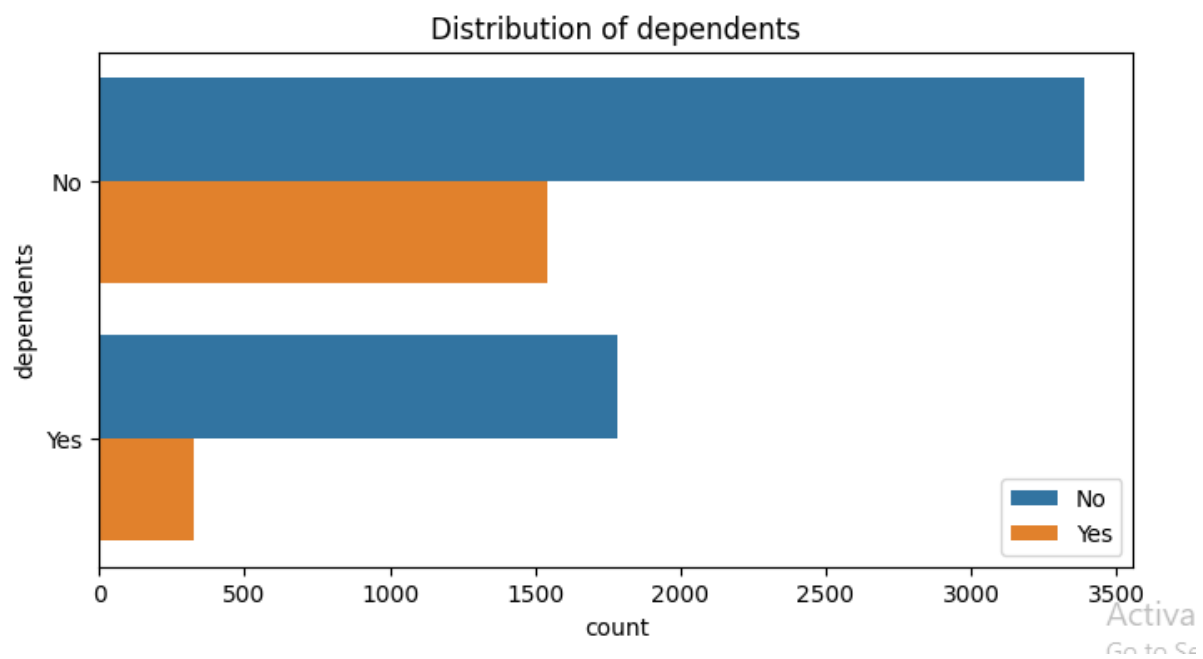
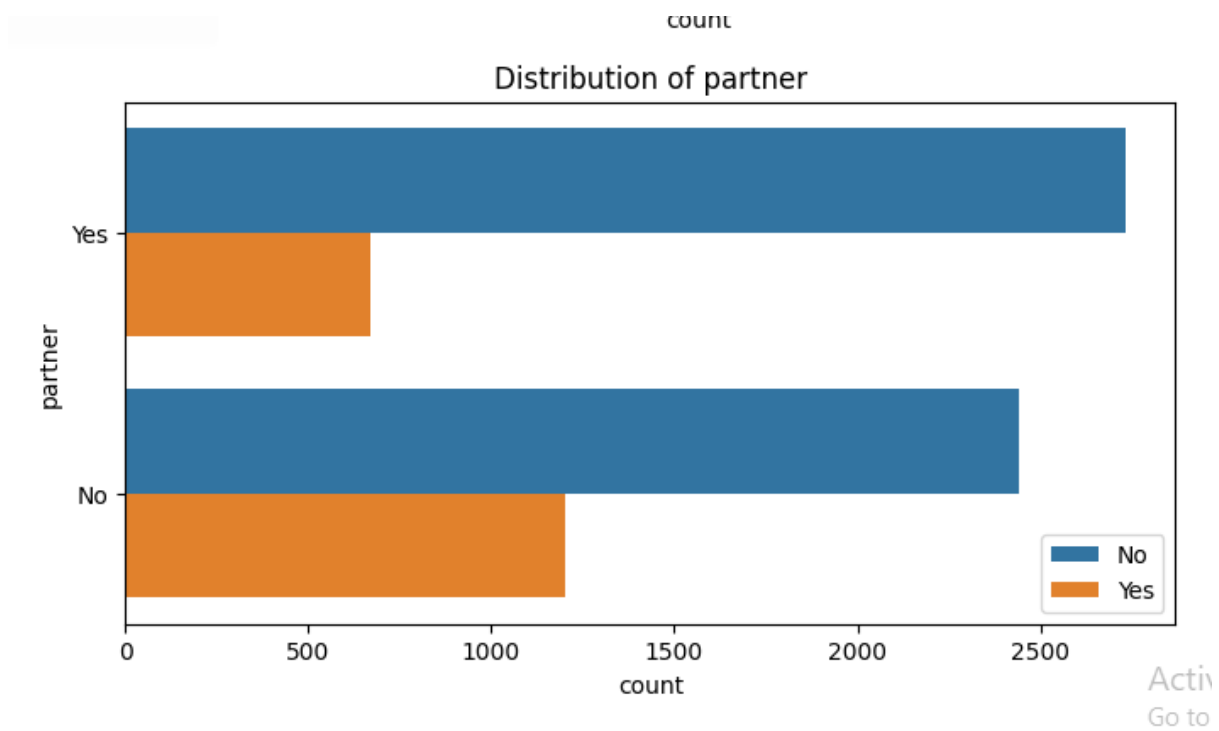
3. Exploratory Data Analysis (EDA)

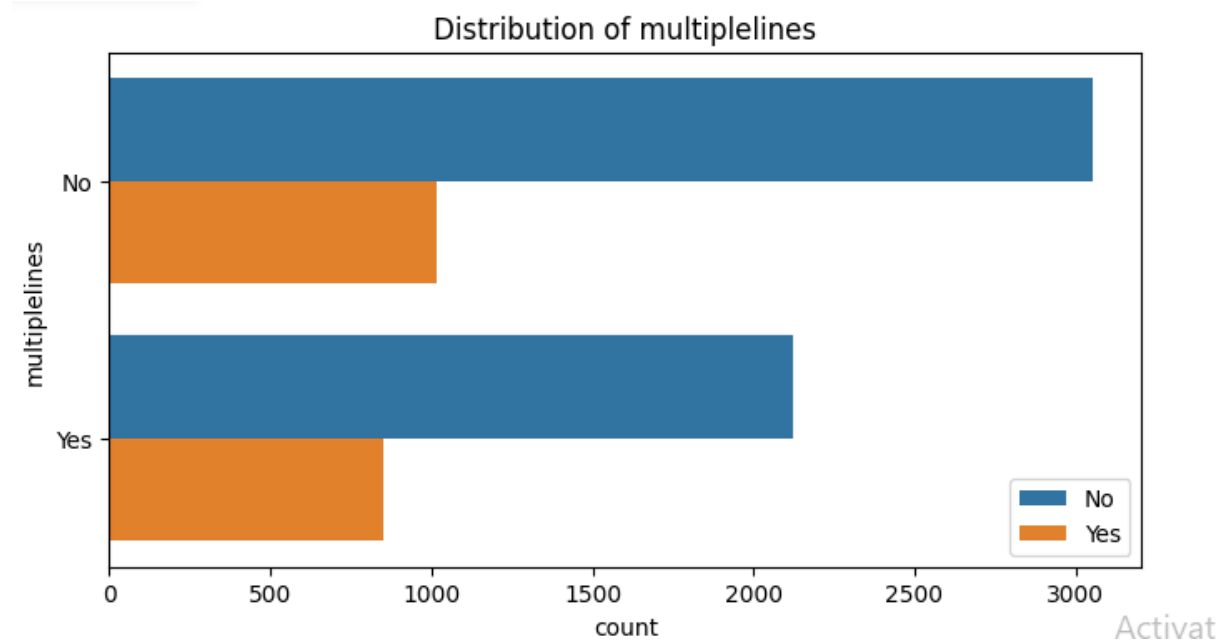
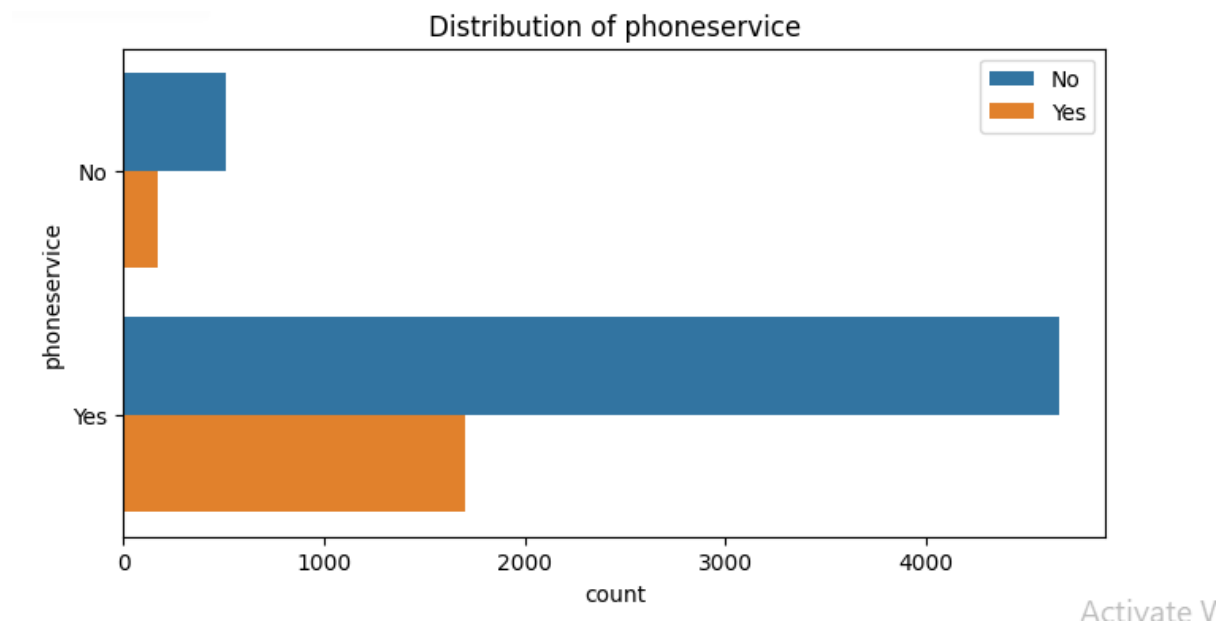
3.1 Visualization :

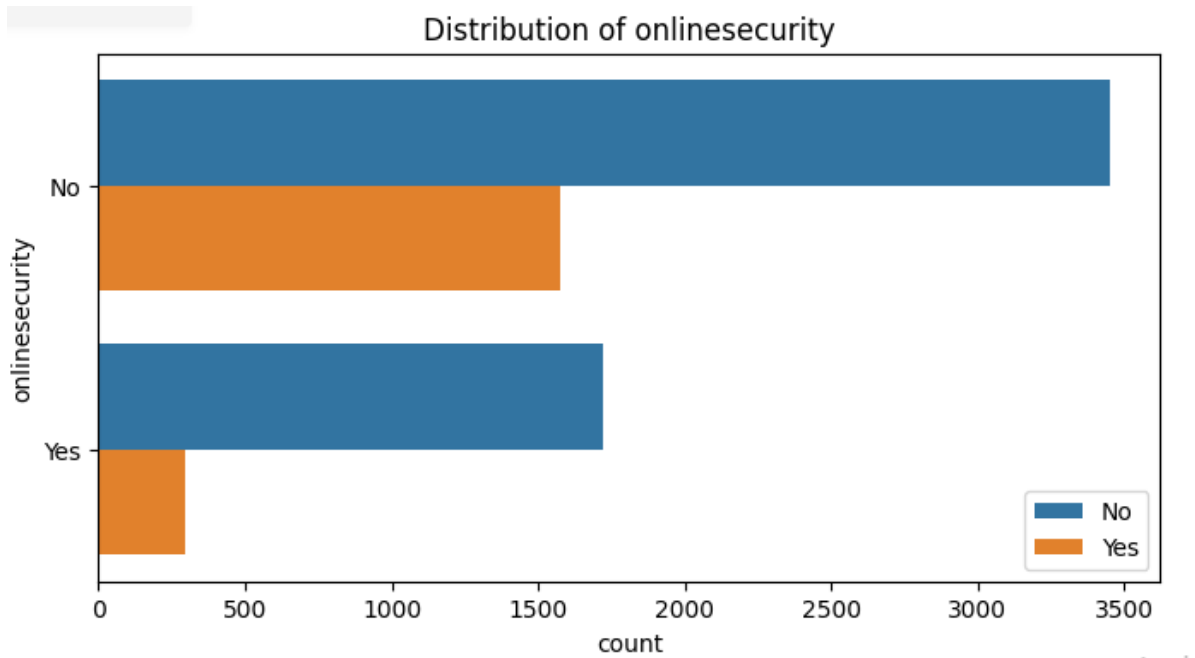
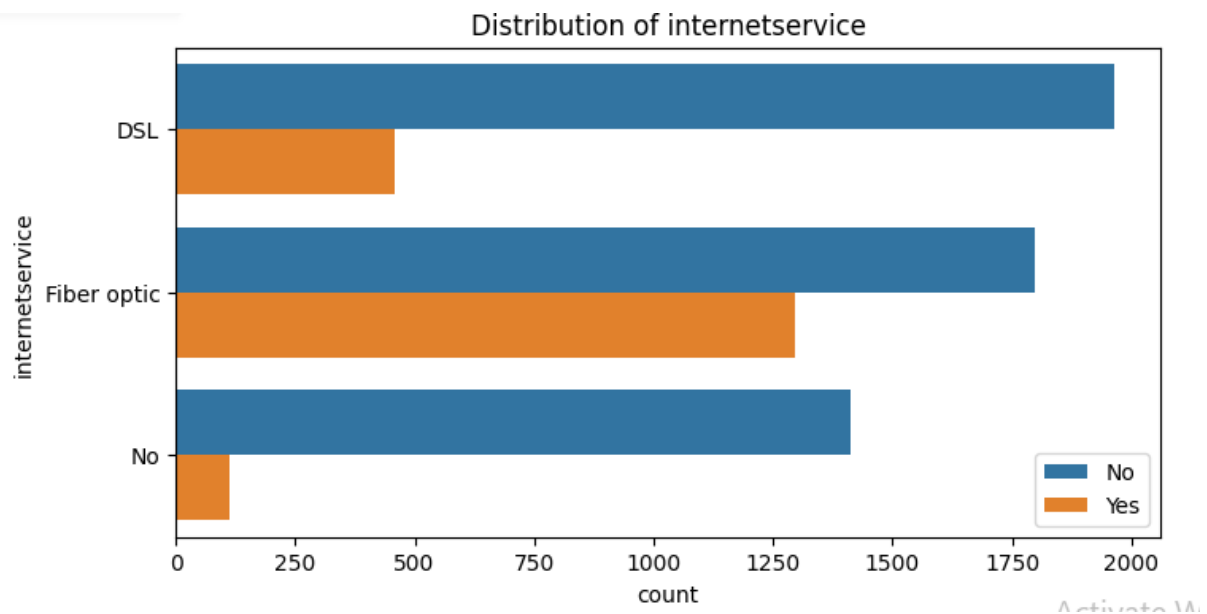
- Categorical Features: Count plots to visualize the distribution and relationship with churn.

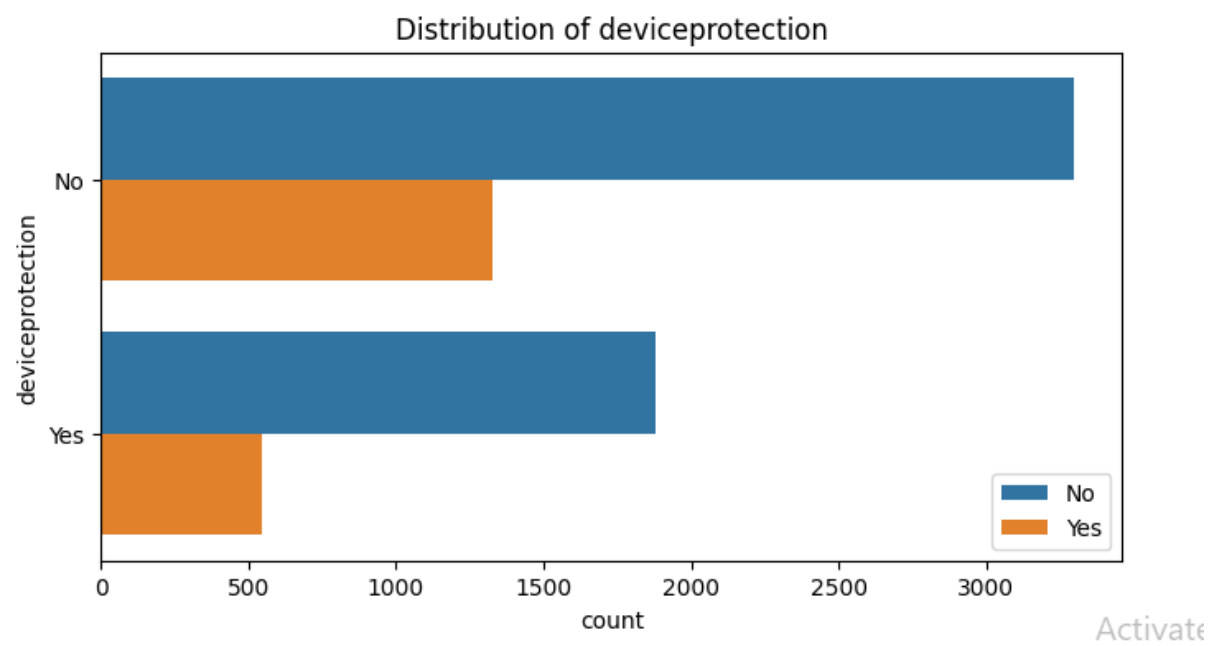
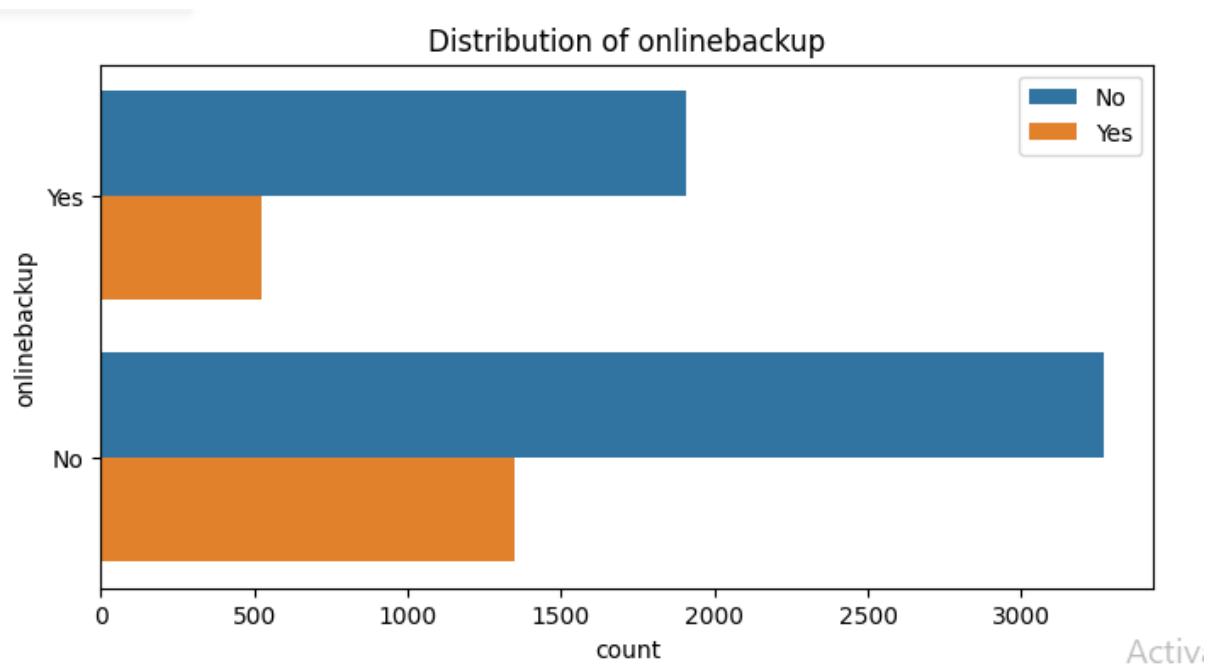
```
for i in categorical:
    plt.figure(figsize=(8, 4))
    sns.countplot(y=i, data=churn_data, hue='churn')
    plt.title(f'Distribution of {i}')
    plt.legend()
    plt.show()
```

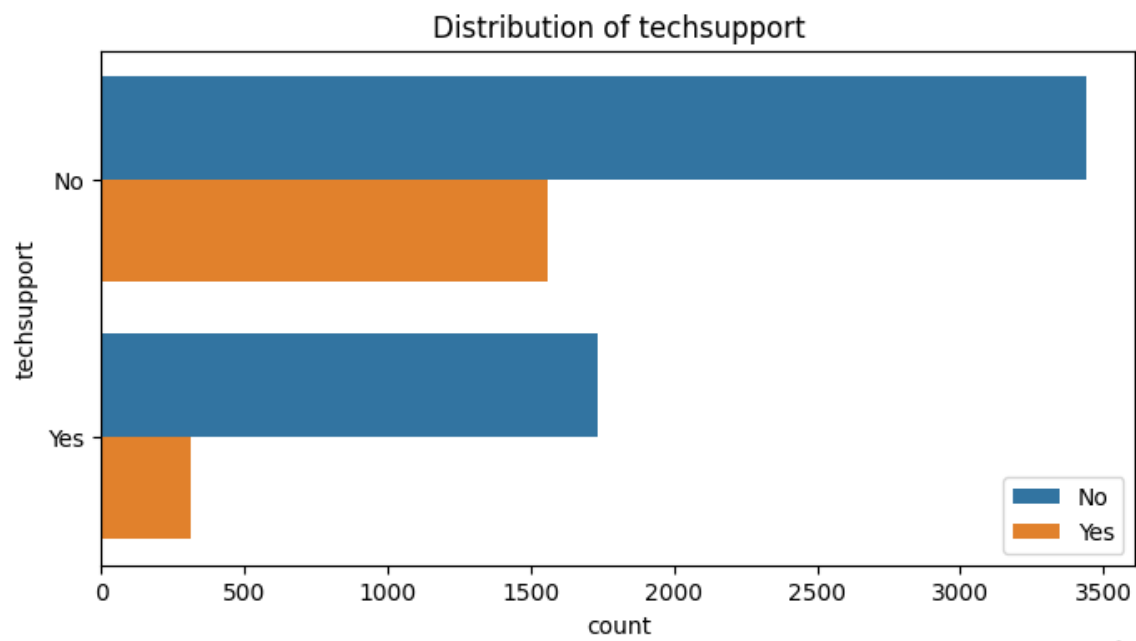




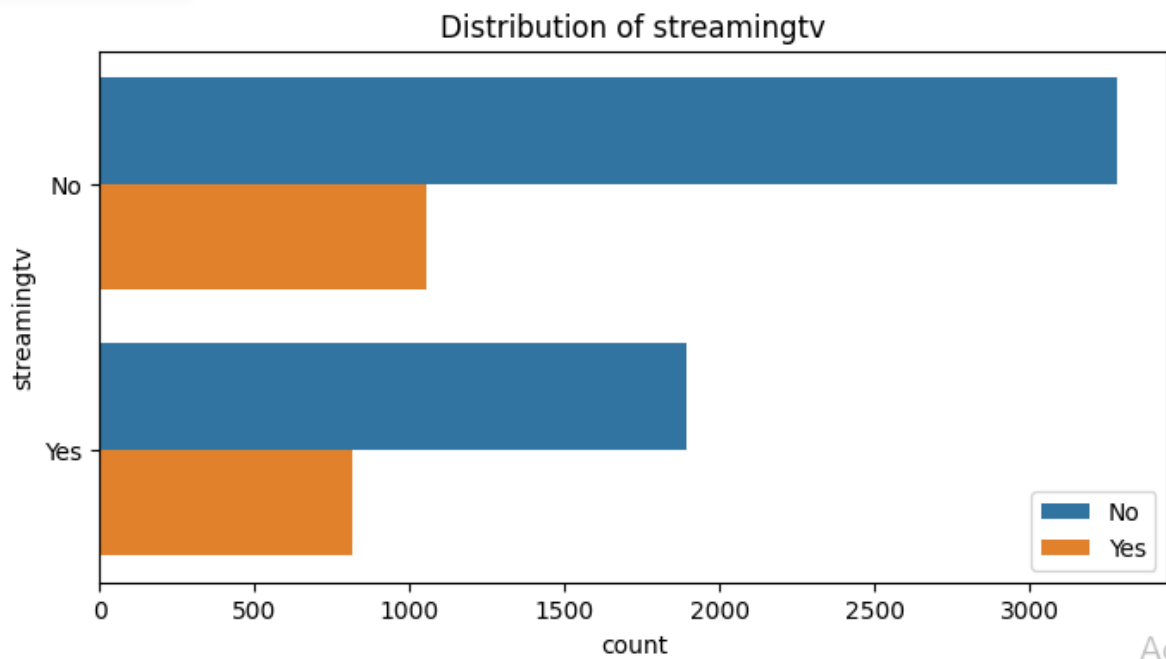





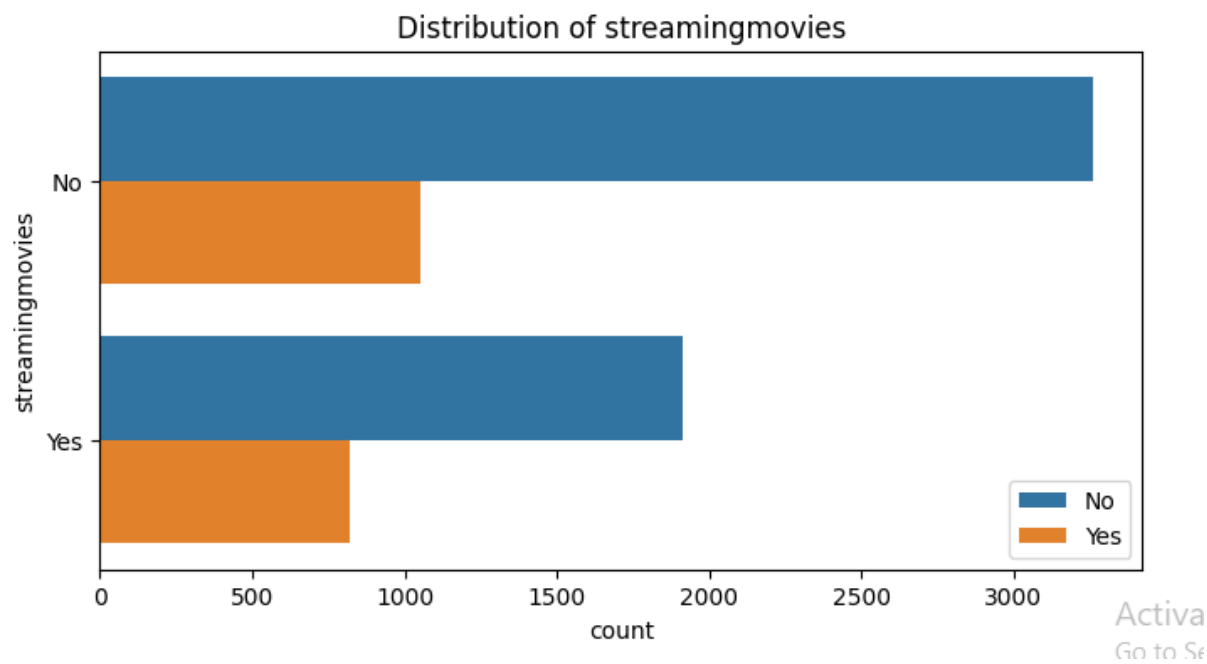



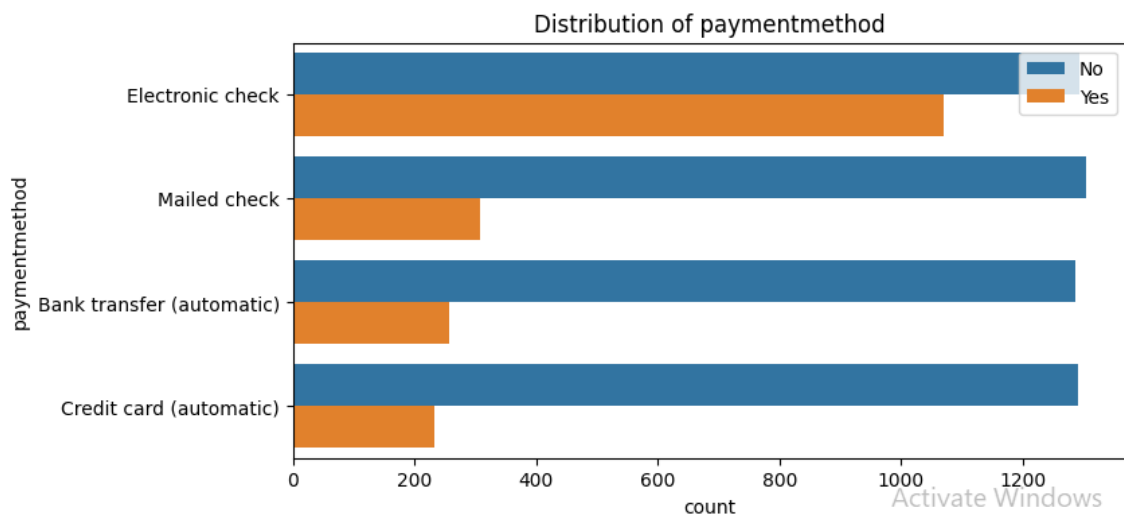
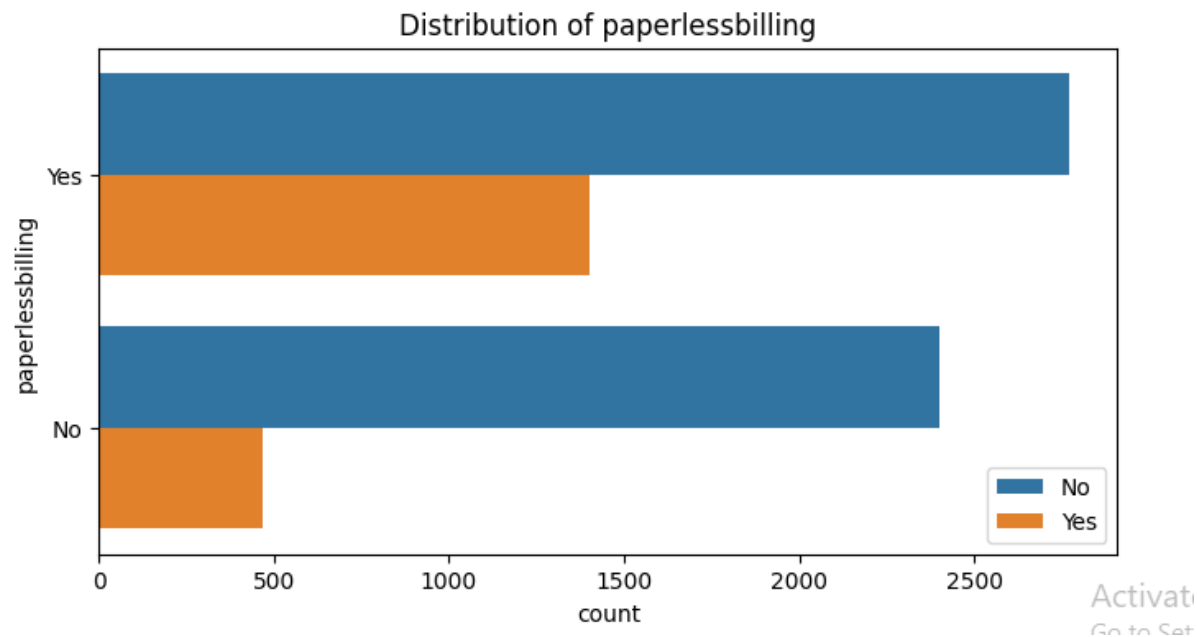


Active




Active





4. Data Balancing

```
import plotly.express as px
fig = px.pie(churn_data, names="churn", hole = 0.4, template = "plotly_dark")
fig.show()
```



- Technique Used: SMOTE (Synthetic Minority Oversampling Technique)

```
x = churn_data.drop("churn",axis=1)
y = churn_data["churn"]
```

```
from imblearn.over_sampling import SMOTE
from collections import Counter
print(Counter(y))
```

```
Counter({0: 5174, 1: 1869})
```

```
smote = SMOTE()
x_resampled, y_resampled = smote.fit_resample(x, y)
```




```
x_resampled.shape, y_resampled.shape
```

```
((10348, 19), (10348,))
```

```
print(Counter(y_resampled))
```

```
Counter({0: 5174, 1: 5174})
```

```
data = x_resampled  
target = y_resampled
```



- Purpose: Address class imbalance in the churn target variable.

5. Model Training and Evaluation (Without Feature Selection)

5.1 Data Splitting:

- Training and Testing Sets: Description of how the data was split.

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(data, target, test_size=0.2, random_state=42)
```

5.2 Model Training:

- Models Used: Logistic Regression, Decision Tree, Random Forest, KNeighborsClassifier, SVC, XGBoost, GradientBoostingClassifier

LogisticRegression:

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

model_lr = LogisticRegression(max_iter=2000)

model_lr.fit(X_train,y_train)

pred_lr = model_lr.predict(X_test)
pred_train=model_lr.predict(X_train)
accuracy_score_train=accuracy_score(y_train,pred_train)
print('training accuray',accuracy_score_train*100)
accuracy_score_lr = accuracy_score(y_test,pred_lr)
print('testing_accuracy',accuracy_score_lr*100)

sns.heatmap((confusion_matrix(y_test,pred_lr)),annot=True,fmt='.5g',cmap="YlGn");
plt.xlabel('True Values')
plt.ylabel('Predicted Values');

print("")
print("Classification Report")
print(classification_report(y_test, pred_lr))
```

Activat

Decision Tree:

```
from sklearn.tree import DecisionTreeClassifier

for i in range(1,11):
    model_dt = DecisionTreeClassifier(max_depth=i)
    model_dt.fit(X_train,y_train)
    pred_dt = model_dt.predict(X_test)
    accuracy_score_dt = accuracy_score(y_test,pred_dt)
    print(i,accuracy_score_dt)
```

RandomForestClassifier:

```
from sklearn.ensemble import RandomForestClassifier

model_rf = RandomForestClassifier()

model_rf.fit(X_train,y_train)
pred_rf = model_rf.predict(X_test)
pred_train=model_rf.predict(X_train)
print("Predicted: ",Counter(pred_rf))
print("Actual: ",Counter(y_test))

accuracy_score_train=accuracy_score(y_train,pred_train)
print('training accuray',accuracy_score_train*100)
accuracy_score_rf = accuracy_score(y_test,pred_rf)
print('testing_accuracy',accuracy_score_rf*100)

sns.heatmap((confusion_matrix(y_test,pred_rf)),annot=True,fmt='.5g',cmap="YlGn");
plt.xlabel('True Values')
plt.ylabel('Predicted Values');

print(classification_report(y_test, pred_rf))
```

KNN:

```
from sklearn.neighbors import KNeighborsClassifier

model_knn = KNeighborsClassifier()

for i in range(1,30,5):
    model_knn = KNeighborsClassifier(n_neighbors=i)
    model_knn.fit(X_train,y_train)
    pred_knn = model_knn.predict(X_test)
    accuracy_score_knn = accuracy_score(y_test,pred_knn)
    print(i,accuracy_score_knn)
```

Support Vector Classifier:

```
from sklearn.svm import SVC

model_svm = SVC(kernel="rbf")

model_svm.fit(X_train,y_train)

pred_train=model_svm.predict(X_train)
pred_svm = model_svm.predict(X_test)

accuracy_score_train=accuracy_score(y_train,pred_train)
print('training accuray',accuracy_score_train*100)
accuracy_score_svm = accuracy_score(y_test,pred_svm)
print('testing_accuracy',accuracy_score_svm*100)

sns.heatmap((confusion_matrix(y_test,pred_svm)),annot=True,fmt='.5g',cmap="YlGn");
plt.xlabel('True Values')
plt.ylabel('Predicted Values');

print(classification_report(y_test, pred_svm))
```

GradientBoostingClassifier:

```
from sklearn.ensemble import GradientBoostingClassifier
model_gbc = GradientBoostingClassifier(n_estimators=100,learning_rate=0.03)

model_gbc.fit(X_train,y_train)
pred_gbc = model_gbc.predict(X_test)
pred_train=model_gbc.predict(X_train)

accuracy_score_train=accuracy_score(y_train,pred_train)
print('training accuray',accuracy_score_train*100)
accuracy_score_gbc = accuracy_score(y_test,pred_gbc)
print('testing_accuracy',accuracy_score_gbc*100)

sns.heatmap((confusion_matrix(y_test,pred_gbc)),annot=True,fmt='.5g',cmap="YlGn");
plt.xlabel('True Values')
plt.ylabel('Predicted Values');
```

XGBClassifier:

```
from xgboost import XGBClassifier

model_xgb = XGBClassifier(n_estimators=100,learning_rate=0.03)

model_xgb.fit(X_train,y_train)
pred_xgb = model_xgb.predict(X_test)
pred_train=model_xgb.predict(X_train)
accuracy_score_train=accuracy_score(y_train,pred_train)
print('training_accuray',accuracy_score_train*100)
accuracy_score_xgb = accuracy_score(y_test,pred_xgb)
print('testing_accuracy',accuracy_score_xgb*100)

sns.heatmap((confusion_matrix(y_test,pred_xgb)),annot=True,fmt='.5g',cmap="YlGn");
plt.xlabel('True Values')
plt.ylabel('Predicted Values');

print(classification_report(y_test, pred_xgb))
```

5.3 Model Evaluation:

- Metrics: Accuracy score, confusion matrix, classification report.


Logistic Regression:

```
training_accuray 81.85552065716357
testing_accuracy 81.78743961352657
```

```
Classification Report
              precision    recall  f1-score   support

     0       0.83         0.80         0.81         1021
     1       0.81         0.84         0.82         1049

 accuracy          0.82         0.82         0.82         2070
 macro avg         0.82         0.82         0.82         2070
weighted avg         0.82         0.82         0.82         2070
```



RandomForestClassifier:

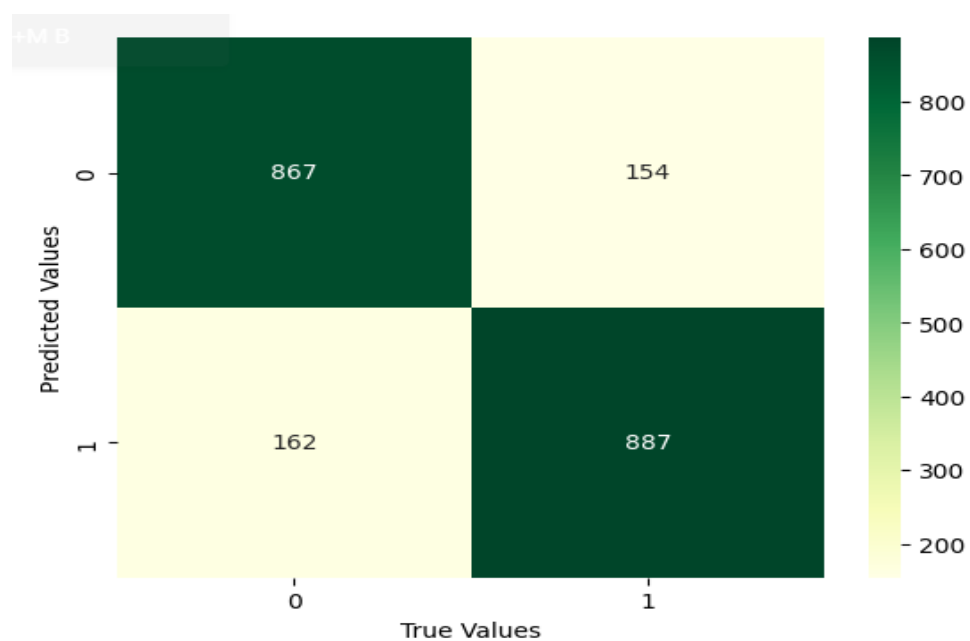
Predicted: Counter({1: 1041, 0: 1029})

Actual: Counter({1: 1049, 0: 1021})

training_accuracy 99.86711766127084

testing_accuracy 84.73429951690822

	precision	recall	f1-score	support
0	0.84	0.85	0.85	1021
1	0.85	0.85	0.85	1049
accuracy			0.85	2070
macro avg	0.85	0.85	0.85	2070
weighted avg	0.85	0.85	0.85	2070




Support Vector Classifier:

```
training_accuracy 65.92172022227591
testing_accuracy 65.26570048309178
      precision    recall  f1-score   support

     0         0.63     0.71     0.67     1021
     1         0.68     0.59     0.63     1049

 accuracy          0.65          2070
  macro avg         0.66         0.65         0.65          2070
 weighted avg         0.66         0.65         0.65          2070
```

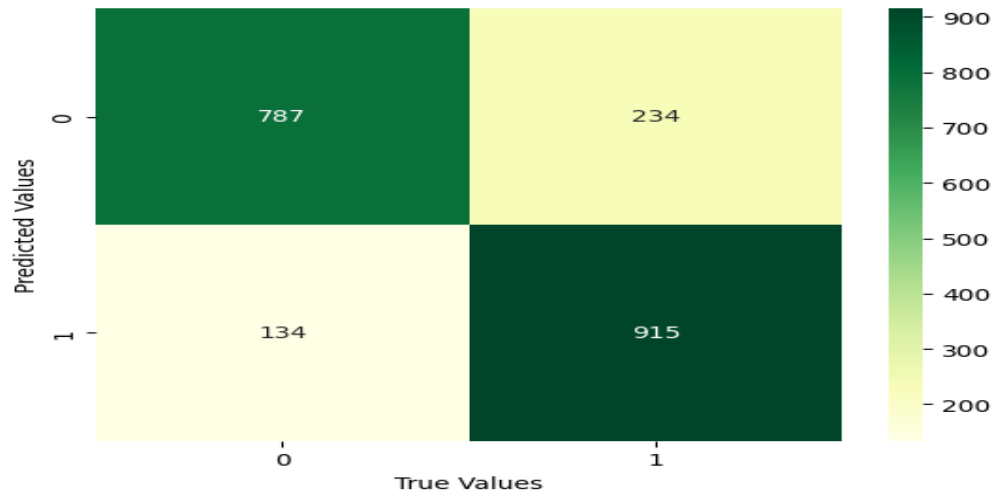


GradientBoostingClassifier:

```
training_accuracy 82.08504469678667
testing_accuracy 82.22222222222221
      precision    recall  f1-score   support

     0         0.85     0.77     0.81     1021
     1         0.80     0.87     0.83     1049

 accuracy          0.82          2070
  macro avg         0.83         0.82         0.82          2070
 weighted avg         0.83         0.82         0.82          2070
```




XGBClassifier:

training_accuracy 85.92655230732062

testing_accuracy 83.04347826086956

	precision	recall	f1-score	support
0	0.86	0.79	0.82	1021
1	0.81	0.87	0.84	1049
accuracy			0.83	2070
macro avg	0.83	0.83	0.83	2070
weighted avg	0.83	0.83	0.83	2070



```
models = pd.DataFrame({
    "Model": ["Logistic Regression",
              "Decision Tree",
              "Random Forest",
              "KNN",
              "SVM",
              "XGBoost",
              "GradientBoosting"],

    "Accuracy Score" : [accuracy_score_lr*100,accuracy_score_dt*100,accuracy_score_rf*100,accuracy_score_knn*100,
                        accuracy_score_svm*100,accuracy_score_xgb*100,accuracy_score_gbc*100]
})

print("Model Accuracies without using RFE")
models
```

Model Accuracies without using RFE

	Model	Accuracy Score
0	Logistic Regression	81.787440
1	Decision Tree	80.966184
2	Random Forest	84.734300
3	KNN	76.183575
4	SVM	65.265700
5	XGBoost	83.043478
6	GradientBoosting	82.222222

6. Feature Selection using RFE

6.1 Recursive Feature Elimination (RFE)

- Model Used: Logistic Regression

```
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression
```


```
scores = []

# Perform RFE with cross-validation to determine the optimal number of features
for n_features in range(1, X_train.shape[1] + 1):
    rfe = RFE(estimator=model_lr, n_features_to_select=n_features)
    rfe.fit(X_train, y_train)
    score = cross_val_score(rfe, X_train, y_train, cv=5, scoring='accuracy').mean()
    scores.append(score)
```



```
plt.figure(figsize=(10, 6))
plt.plot(range(1, X_train.shape[1] + 1), scores, marker='o')
plt.xlabel('Number of Features Selected')
plt.ylabel('Cross-Validated Accuracy')
plt.title('Optimal Number of Features - RFE')
plt.show()

# Get the optimal number of features
optimal_n_features = np.argmax(scores) + 1
print(f'Optimal number of features: {optimal_n_features}')
```



- Objective: Identify optimal number of features

```
lg=LogisticRegression(max_iter=2000)
lg.fit(X_train,y_train)
rfe=RFE(estimator=lg,n_features_to_select=16)
rfe=rfe.fit(X_train,y_train)
```

```
selected_features = X_train.columns[rfe.support_]
print("Selected features:", selected_features)
```

```
Selected features: Index(['gender', 'seniorcitizen', 'partner', 'dependents', 'phoneservice',
                        'multiplelines', 'internetservice', 'onlinesecurity', 'onlinebackup',
                        'deviceprotection', 'techsupport', 'streamingtv', 'streamingmovies',
                        'contract', 'paperlessbilling', 'monthlycharges'],
                        dtype='object')
```

```
import statsmodels.api as sm
X_train_sm=sm.add_constant(X_train_rfe)
logistic_model=sm.Logit(y_train,X_train_rfe).fit()
logistic_model.summary()
```

Optimization terminated successfully.

Current function value: 0.417370

Iterations 7

Logit Regression Results

Dep. Variable:	churn	No. Observations:	8278
Model:	Logit	Df Residuals:	8262
Method:	MLE	Df Model:	15
Date:	Mon, 10 Jun 2024	Pseudo R-squ.:	0.3979
Time:	13:30:15	Log-Likelihood:	-3455.0
converged:	True	LL-Null:	-5737.8
Covariance Type:	nonrobust	LLR p-value:	0.000

	coef	std err	z	P> z	[0.025	0.975]
gender	-0.4234	0.059	-7.224	0.000	-0.538	-0.308
seniorcitizen	-0.4733	0.085	-5.547	0.000	-0.641	-0.306
partner	-0.4832	0.073	-6.661	0.000	-0.625	-0.341
dependents	-0.5636	0.089	-6.337	0.000	-0.738	-0.389
phoneservice	-1.5468	0.116	-13.355	0.000	-1.774	-1.320
multiplelines	-0.5331	0.074	-7.191	0.000	-0.678	-0.388
internetservice	-0.0500	0.054	-0.930	0.352	-0.155	0.055
onlinesecurity	-1.2524	0.084	-14.867	0.000	-1.417	-1.087
onlinebackup	-0.9323	0.074	-12.595	0.000	-1.077	-0.787
deviceprotection	-0.6995	0.078	-8.992	0.000	-0.852	-0.547
techsupport	-1.0538	0.085	-12.404	0.000	-1.220	-0.887
streamingtv	-0.5319	0.083	-6.418	0.000	-0.694	-0.369
streamingmovies	-0.4982	0.083	-6.020	0.000	-0.660	-0.336
contract	-1.1368	0.064	-17.855	0.000	-1.262	-1.012
paperlessbilling	-0.1264	0.065	-1.931	0.053	-0.255	0.002
monthlycharges	0.0573	0.002	29.931	0.000	0.054	0.061

```
X_train_rfe=X_train_rfe.drop(['paperlessbilling','internetservice'],axis=1)
```

```
X_test_rfe=X_test[X_train_rfe.columns]
```

6.2 Feature Selection

- Best Features: List of selected features based on RFE analysis.

	coef	std err	z	P> z	[0.025	0.975]
gender	-0.4339	0.058	-7.435	0.000	-0.548	-0.320
seniorcitizen	-0.4914	0.085	-5.780	0.000	-0.658	-0.325
partner	-0.4892	0.072	-6.756	0.000	-0.631	-0.347
dependents	-0.5655	0.089	-6.364	0.000	-0.740	-0.391
phoneservice	-1.6123	0.096	-16.870	0.000	-1.800	-1.425
multiplelines	-0.5338	0.074	-7.230	0.000	-0.678	-0.389
onlinesecurity	-1.2360	0.084	-14.749	0.000	-1.400	-1.072
onlinebackup	-0.9326	0.074	-12.578	0.000	-1.078	-0.787
deviceprotection	-0.6950	0.078	-8.925	0.000	-0.848	-0.542
techsupport	-1.0451	0.085	-12.360	0.000	-1.211	-0.879
streamingtv	-0.5398	0.083	-6.527	0.000	-0.702	-0.378
streamingmovies	-0.4987	0.083	-6.028	0.000	-0.661	-0.337
contract	-1.1398	0.063	-18.131	0.000	-1.263	-1.017
monthlycharges	0.0566	0.002	31.557	0.000	0.053	0.060

7. Model Training and Evaluation (With Feature Selection)

7.1 Retraining:

- Model Used: Logistic Regression with selected features.

```
logreg = LogisticRegression()

# Train the classifier
logreg.fit(X_train_rfe, y_train)
```

7.2 Model Evaluation:

- Metrics: Accuracy score, precision, recall, F1 score, confusion matrix, classification report.

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report

y_pred_logreg = logreg.predict(X_test_rfe)
accuracy_logreg = accuracy_score(y_test, y_pred_logreg) * 100
conf_matrix_logreg = confusion_matrix(y_test, y_pred_logreg)
class_report_logreg = classification_report(y_test, y_pred_logreg)
precision_logreg = precision_score(y_test, y_pred_logreg) * 100
recall_logreg = recall_score(y_test, y_pred_logreg) * 100
f1_logreg = f1_score(y_test, y_pred_logreg) * 100
```

```
print("Logistic Regression Classifier Metrics:")
print("Accuracy:", accuracy_logreg)
print("Precision:", precision_logreg)
print("Recall:", recall_logreg)
print("F1 Score:", f1_logreg)
```

```
Logistic Regression Classifier Metrics:
Accuracy: 82.41545893719807
Precision: 80.82808280828083
Recall: 85.60533841754051
F1 Score: 83.14814814814815
```

8. Hyperparameter Tuning using GridSearchCV

8.1 Pipelines and Parameter Grids:

- Models Included: Logistic Regression, Random Forest, Gradient Boosting, Decision Tree

```
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split, GridSearchCV
```

Pipelines:

```
# Logistic Regression Pipeline
pipeline_lr = Pipeline([
    ('scaler', StandardScaler()),
    ('classifier', LogisticRegression(max_iter=2000))
])

# Random Forest Pipeline
pipeline_rf = Pipeline([
    ('classifier', RandomForestClassifier())
])

# Gradient Boosting Pipeline
pipeline_gb = Pipeline([
    ('classifier', GradientBoostingClassifier())
])

# Decision Tree Pipeline
pipeline_dt = Pipeline([
    ('classifier', DecisionTreeClassifier())
])
```

```

# Define parameter grids for each model
param_grid_lr = {
    'classifier__solver': ['liblinear', 'saga'],
    'classifier__C': [0.1, 1, 10],
    'classifier__penalty': ['l2']
}

param_grid_rf = {
    'classifier__n_estimators': [10,100,200,300],
    'classifier__max_depth': [None, 10, 20],
    'classifier__min_samples_split': [2, 5],
    'classifier__min_samples_leaf': [1, 2]
}

param_grid_gb = {
    'classifier__n_estimators': [10,100,200,300],
    'classifier__learning_rate': [0.1, 0.01],
    'classifier__max_depth': [1,2,3,4,6,8]
}

param_grid_dt = {
    'classifier__max_depth': [None, 10, 20],
    'classifier__min_samples_split': [2, 5],
    'classifier__min_samples_leaf': [1, 2]
}

```

```

pipelines = {
    'Logistic Regression': (pipeline_lr, param_grid_lr),
    'Random Forest': (pipeline_rf, param_grid_rf),
    'Gradient Boosting': (pipeline_gb, param_grid_gb),
    'Decision Tree': (pipeline_dt, param_grid_dt)
}

```

- Objective: Tune hyperparameters to maximize accuracy.

8.2 GridSearchCV:

- Process: Description of GridSearchCV process.

```

best_estimators = {}
for name, (pipeline, param_grid) in pipelines.items():
    grid_search = GridSearchCV(pipeline, param_grid, cv=5, n_jobs=-1, scoring='accuracy')
    grid_search.fit(X_train_rfe, y_train)
    best_estimators[name] = grid_search.best_estimator_
    print(f"Best parameters for {name}: {grid_search.best_params_}")
    print(f"Best cross-validation accuracy for {name}: {grid_search.best_score_}")

```

- Best Hyperparameters: Summary of the best hyperparameter combinations found.

9. Final Model Evaluation

9.1 Evaluation Metrics:

- Metrics: Accuracy score, confusion matrix, classification report.

```
Logistic Regression
Train_Accuracy: 0.8061125875815415
Test Accuracy: 0.8236714975845411
Confusion Matrix:
[[807 214]
 [151 898]]
Classification Report:
              precision    recall  f1-score   support

     0           0.84       0.79       0.82       1021
     1           0.81       0.86       0.83       1049

 accuracy          0.82
 macro avg         0.82
 weighted avg      0.82
```

```
Random Forest
Train_Accuracy: 0.858661512442619
Test Accuracy: 0.8280193236714976
Confusion Matrix:
[[806 215]
 [141 908]]
Classification Report:
              precision    recall  f1-score   support

     0           0.85       0.79       0.82       1021
     1           0.81       0.87       0.84       1049

 accuracy          0.83
 macro avg         0.83
 weighted avg      0.83
```

Gradient Boosting

Train_Accuracy: 0.8235080937424498

Test Accuracy: 0.8265700483091788

Confusion Matrix:

[[809 212]

[147 902]]

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.79	0.82	1021
1	0.81	0.86	0.83	1049
accuracy			0.83	2070
macro avg	0.83	0.83	0.83	2070
weighted avg	0.83	0.83	0.83	2070

Decision Tree

Train_Accuracy: 0.8518965933800435

Test Accuracy: 0.8028985507246377

Confusion Matrix:

[[776 245]



[163 886]]



Classification Report:

	precision	recall	f1-score	support
0	0.83	0.76	0.79	1021
1	0.78	0.84	0.81	1049
accuracy			0.80	2070
macro avg	0.80	0.80	0.80	2070
weighted avg	0.80	0.80	0.80	2070

9.2 Confusion Matrix Visualization:

- Purpose: Understand model performance on different classes.





10. Conclusion

Summary: Comprehensive comparison of models with and without feature selection and hyperparameter tuning.

without feature selection:

Logistic Regression

Train_Accuracy: 0.8136023194008215

Test Accuracy: 0.8251207729468599

Confusion Matrix:

[[809 212]

[150 899]]

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.79	0.82	1021
1	0.81	0.86	0.83	1049
accuracy			0.83	2070
macro avg	0.83	0.82	0.82	2070
weighted avg	0.83	0.83	0.82	2070

Random Forest

Train_Accuracy: 0.998550374486591

Test Accuracy: 0.8492753623188406

Confusion Matrix:

[[853 168]

[144 905]]

Classification Report:

	precision	recall	f1-score	support
0	0.86	0.84	0.85	1021
1	0.84	0.86	0.85	1049
accuracy			0.85	2070
macro avg	0.85	0.85	0.85	2070
weighted avg	0.85	0.85	0.85	2070

Gradient Boosting

Train_Accuracy: 0.9981879681082387

Test Accuracy: 0.8545893719806763

Confusion Matrix:

```
[[874 147]
```

```
[154 895]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.86	0.85	1021
1	0.86	0.85	0.86	1049
accuracy			0.85	2070
macro avg	0.85	0.85	0.85	2070
weighted avg	0.85	0.85	0.85	2070

Decision Tree

Train_Accuracy: 0.8706209229282436

Test Accuracy: 0.8193236714975846

Confusion Matrix:

```
[[780 241]
```

```
[133 916]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.76	0.81	1021
1	0.79	0.87	0.83	1049
accuracy			0.82	2070
macro avg	0.82	0.82	0.82	2070
weighted avg	0.82	0.82	0.82	2070

Without feature selection Gradient Boosting is better than the other algorithms. But is Overfitting .

With feature selection:

Logistic Regression

Train_Accuracy: 0.8061125875815415

Test Accuracy: 0.8236714975845411

Confusion Matrix:

[[807 214]

[151 898]]

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.79	0.82	1021
1	0.81	0.86	0.83	1049
accuracy			0.82	2070
macro avg	0.82	0.82	0.82	2070
weighted avg	0.82	0.82	0.82	2070

Random Forest

Train_Accuracy: 0.858661512442619

Test Accuracy: 0.8280193236714976

Confusion Matrix:

[[806 215]

[141 908]]

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.79	0.82	1021
1	0.81	0.87	0.84	1049
accuracy			0.83	2070
macro avg	0.83	0.83	0.83	2070
weighted avg	0.83	0.83	0.83	2070

Gradient Boosting

Train_Accuracy: 0.8235080937424498

Test Accuracy: 0.8265700483091788

Confusion Matrix:

[[809 212]

[147 902]]

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.79	0.82	1021
1	0.81	0.86	0.83	1049
accuracy			0.83	2070
macro avg	0.83	0.83	0.83	2070
weighted avg	0.83	0.83	0.83	2070

```

Decision Tree
Train_Accuracy: 0.8518965933800435
Test Accuracy: 0.8028985507246377
Confusion Matrix:
[[776 245]
 [163 886]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.83	0.76	0.79	1021
1	0.78	0.84	0.81	1049
accuracy			0.80	2070
macro avg	0.80	0.80	0.80	2070
weighted avg	0.80	0.80	0.80	2070

Without feature selection Gradient Boosting and Random Forest is better than the other algorithms. But is Random Forest is less Overfitting when compare with without feature selection .The best model is Gradient Boosting because is not overfitted.

- Best Performing Model:Not explicitly mentioned, but comparative analysis provided.